

Class 10: Structural Bioinformatics 1

Christeena Pano (PID: A17842497)

2026-04-30

Table of contents

Background	1
Introduction to the RCSB Protein Data Bank (PDB)	1
PDB Statistics and Data Import	2
Visualizing PDB data with Mol-star	5
Getting started with the Bio3D package	10
Predict the flexibility of a given structure	12
Comparative analysis of the ADK family	13
Principal Component Analysis	16

Background

The main repository of high-resolution structural data on biomolecules is called the **Protein Data Bank** (PDB).

Introduction to the RCSB Protein Data Bank (PDB)

The PDB archive is the major repository of information about the 3D structures of large biological molecules, including proteins and nucleic acids. Understanding the shape of these molecules helps to understand how they work. This knowledge can be used to help deduce a structure's role in human health and disease, and in drug development. The structures in the PDB range from tiny proteins and bits of DNA or RNA to complex molecular machines like the ribosome composed of many chains of protein and RNA.

In the first section of this lab we will interact with the main US based PDB website.

PDB Statistics and Data Import

What is in the PDB in terms of molecule type and structure determination method.

Read a CSV file of current PDB stats obtained from <https://www.rcsb.org/stats/summary>

```
# View data using read.csv()
pdb <- read.csv("Data Export Summary.csv")
pdb
```

	Molecular.Type	X.ray	EM	NMR	Integrative	Multiple.methods
1	Protein (only)	180,758	23,111	12,813	348	229
2	Protein/Oligosaccharide	10,488	3,741	34	8	11
3	Protein/NA	9,205	6,751	287	26	8
4	Nucleic acid (only)	3,154	250	1,578	3	15
5	Other	178	27	35	4	0
6	Oligosaccharide (only)	11	0	6	0	1
	Neutron	Other	Total			
1	84	32	217,375			
2	1	0	14,283			
3	0	0	16,277			
4	3	1	5,004			
5	0	0	244			
6	0	4	22			

Q1: What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy.

The percentage of structures in the PDB solved by X-Ray is 80.48577%. The percentage of structures in the PDB that are solved by Electron Microscopy is 13.38046%.

```
pdb$X.ray
```

```
[1] "180,758" "10,488" "9,205" "3,154" "178" "11"
```

This print out above `pdb$X.ray` is “character” not “numeric”. Therefore I can’t do math with it. We need to fix this...

Two functions that can help here are `sub()` and `as.numeric()`

```
# We want to get rid (or sub out) commas:
x <- pdb$X.ray
tmp <- sub(",", "", x)
sum(as.numeric(tmp))
```

```
[1] 203794
```

We could make a function to do this:

```
# Sub out commas and assign function to rm.comma()
rm.comma <- function(x) {
  tmp <- sub(",", "", x)
  sum(as.numeric(tmp))
}
```

```
# use rm.comma() to find the proportions of structures found by each respective method
n.tot <- rm.comma(pdb$Total)
n.xray <- rm.comma(pdb$X.ray)
n.em <- rm.comma(pdb$EM)

n.xray/n.tot *100
```

```
[1] 80.48577
```

```
n.em/n.tot * 100
```

```
[1] 13.38046
```

We could also use a different import function 'for this CSV that speaks American (i.e. deals with commas in numbers in a comma separated value file)

```
library(readr)

read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 9
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (1): Molecular Type
```

```
dbl (4): Integrative, Multiple methods, Neutron, Other
```

```
num (4): X-ray, EM, NMR, Total
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
# A tibble: 6 x 9
  `Molecular Type`      `X-ray`      EM      NMR Integrative `Multiple methods` Neutron
  <chr>                <dbl> <dbl> <dbl>          <dbl>          <dbl> <dbl>
1 Protein (only)      180758 23111 12813          348          229    84
2 Protein/Oligosacch~  10488  3741   34           8           11     1
3 Protein/NA          9205  6751  287          26           8     0
4 Nucleic acid (only)  3154   250 1578           3          15     3
5 Other               178    27   35           4           0     0
6 Oligosaccharide (o~   11     0    6           0           1     0
# i 2 more variables: Other <dbl>, Total <dbl>
```

Q2: What proportion of structures in the PDB are protein?

The proportion of structures in the PDB that are protein is 97.92%.

```
pdb$Total[1]
```

```
[1] "217,375"
```

The total number of protein sequences is 202,556,314, according to Uniprot database.

```
(217375/202556314) *100
```

```
[1] 0.1073158
```

Key-point: We have a very, very small structural coverage of known proteins (~0.1%). Most structures we know about (~80%) come from one method (X-ray crystallography)

```
# Find proportion of protein structures in the PDB database by using data from rows 1-3 and c
n.tot <- rm.comma(pdb$Total)
n.pro <- rm.comma(pdb$Total[1:3])

(n.pro / n.tot)
```

```
[1] 0.9791868
```

Q3: Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB?

There are 1227 HIV-1 protease structures in the current PDB.

Visualizing PDB data with Mol-star

Main stand alone web version with all features is at <https://molstar.org/viewer/>.




Figure 1: The HIV-protease enzyme is a homodimer with two chains



Figure 2: The HIV-protease enzyme is a homodimer with two chains




Figure 3: The HIV-protease enzyme with both Asp 25 positions selected




Figure 4: HIV-Pr molecular surface representation display of polymer showing the binding cleft site for the inhibitor (drug) molecule

Q4: Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?

Since hydrogen only has one electron, it is not visible with the methods used to solve/find these structures. Electron microscopy and X-ray diffraction better detect the oxygen molecule. Therefore, only one atom (oxygen) is used to represent the water molecules in this structure.

Q5: There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have

This water molecule is HOH 308.

Q6: Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic residues ASP 25 in each chain and the critical water (we recommend “Ball & Stick” for these side-chains). Add this figure to your Quarto document.




Figure 5: HIV-protease cartoon polymer with both chains A and B, the ligand, both Asp 25 residues selected, and the critical water molecule

```
# Pak can install from multiple sources
#install.packages("pak")

#| eval: !expr knitr::is_html_output()

pak::pkg_install(c("bioboot/bio3dview",
                   "NGLVieweR",
                   "bioc::msa"))
```

Loading metadata database

Loading metadata database ... done

No downloads are needed

3 pkgs + 47 deps: kept 44 [4.9s]

```
#install.packages("bio3d")  
library(bio3d)
```

Getting started with the Bio3D package

Bio3D is an R package from CRAN for structural bioinformatics

```
library(bio3d)  
  
pdb <- read.pdb("1hsg")  
pdb
```

Q7: How many amino acid residues are there in this pdb object?

There are 198 amino acid residues in this pdb object.

Q8: Name one of the two non-protein residues?

MK1 is one of the two non-protein residues.

Q9: How many protein chains are in this structure?

There are 2 protein chains in this structure.

```
head(pdb$atom)
```

NULL

There are lots of functions that can work with these **pdb** objects:

```
#head( pdbseq(pdb))
```

```
#| eval: !expr knitr::is_html_output()

library(bio3dview)
library(NGLVieweR)

#| eval: !expr knitr::is_html_output()
#view.pdb(pdb) |> setSpin()
```

```
#install.packages("remotes")
remotes::install_github("bioboot/bio3dview")
```

Using GitHub PAT from the git credential store.

Skipping install of 'bio3dview' from a github remote, the SHA1 (568c6e99) has not changed since last install. Use `force = TRUE` to force installation

```
#install.packages("NGLVieweR")
```

let's try a custom view

```
#view.pdb(pdb,colorScheme="sse", backgroundColor="black")
```

Q. Create a custom view highlighting the active site ASP (resno25), the two chains (in your choice of colors) and the ligand all on a custom color background

```
library(NGLVieweR)
active.site <- atom.select(pdb, resno=25)

#| eval: !expr knitr::is_html_output()

#view.pdb(pdb,
  #cols= c("red", "blue"),
  #highlight = active.site,
  #highlight.style = "spacefill",
  #backgroundColor= "pink") |>
#setRock()
```

Predict the flexibility of a given structure

Let's do a Normal Mode Analysis (NMA) to predict the flexibility of a given `pdb` object:

```
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file
PDB has ALT records, taking A only, `rm.alt=TRUE`

```
adk
```

```
Call: read.pdb(file = "6s36")
```

```
Total Models#: 1
  Total Atoms#: 1898,  XYZs#: 5694  Chains#: 1  (values: A)

  Protein Atoms#: 1654  (residues/Calpha atoms#: 214)
  Nucleic acid Atoms#: 0  (residues/phosphate atoms#: 0)

  Non-protein/nucleic Atoms#: 244  (residues: 244)
  Non-protein/nucleic resid values: [ CL (3), HOH (238), MG (2), NA (1) ]
```

Protein sequence:


```
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMLRAAVKSGSELGKQAKDIMDAGKLV
DELVIALVKERIAQEDCRNGFLLDGFPRTIPQADAMKEAGINVDYVLEFDVPDELIVDKI
VGRRVHAPSGRVYHVKFNPVKVEGKDDVTGEELTTRKDDQEETVRKRLVEYHQMTAPLIG
YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
```

```
+ attr: atom, xyz, seqres, helix, sheet,
      calpha, remark, call
```

```
# Plot data
m <- nma( adk )
```

```
Building Hessian...      Done in 0.015 seconds.
Diagonalizing Hessian... Done in 0.283 seconds.
```

```
plot(m)
```



View the results with an interactive structure view

```
# Create interactive structure using view.nma()
#| eval: !expr knitr::is_html_output()

#view.nma(m)
```

Write out the results for viewing in Mol-star:

```
mktrj(m, file="nma.pdb")
```

Comparative analysis of the ADK family

Q10. Which of the packages above is found only on BioConductor and not CRAN?

msa is found only on BioConductor.

Q11. Which of the above packages is not found on BioConductor or CRAN?

bio3dview is not found on Bioconductor or CRAN.

Q12. True or False? Functions from the pak package can be used to install packages from GitHub and BitBucket?

This is true.

Our first step is find a sequence for this family. We will use the database ID “1ake_A” here:

```
id <- "1ake_A"
aa <- get.seq(id)
```

Warning in get.seq(id): Removing existing file: seqs.fasta

Fetching... Please wait. Done.

```
aa
```

```

      1      .      .      .      .      .      .      60
pdb|1AKE|A  MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMRLRAAVKSGSELGKQAKDIMDAGKLV
      1      .      .      .      .      .      .      60

      61      .      .      .      .      .      .      120
pdb|1AKE|A  DELVIALVKERIAQEDCRNGFLLDGFRTIPQADAMKEAGINVDYVLEFDVPDELIVDRI
      61      .      .      .      .      .      .      120

     121      .      .      .      .      .      .      180
pdb|1AKE|A  VGRRVHAPSGRVYHVKNPPKVEGKDDVTGEELTRKDDQEETVRKRLVEYHQMTPALIG
     121      .      .      .      .      .      .      180

     181      .      .      .      214
pdb|1AKE|A  YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
     181      .      .      .      214
```

Call:

```
read.fasta(file = outfile)
```

Class:

```
fasta
```

Alignment dimensions:

```
1 sequence rows; 214 position columns (214 non-gap, 0 gap)
```

+ attr: id, ali, call

Q13. How many amino acids are in this sequence, i.e. how long is this sequence?

There are 214 amino acids in this sequence.

Search for related sequences in the database

```
blast <- blast.pdb(aa)
```

Searching ... please wait (updates every 5 seconds) RID = 2D04C9NZ014

.....

Reporting 96 hits

```
head(blast$hit.tbl)
```


	queryid	subjectids	identity	alignmentlength	mismatches	gapopens	q.start
1	Query_4143327	1AKE_A	100.000	214	0	0	1
2	Query_4143327	8BQF_A	99.533	214	1	0	1
3	Query_4143327	4X8M_A	99.533	214	1	0	1
4	Query_4143327	6S36_A	99.533	214	1	0	1
5	Query_4143327	9R6U_A	99.533	214	1	0	1
6	Query_4143327	9R71_A	99.533	214	1	0	1

	q.end	s.start	s.end	evalue	bitscore	positives	mlog.evalue	pdb.id	acc
1	214	1	214	1.82e-156	432	100.00	358.6044	1AKE_A	1AKE_A
2	214	21	234	2.98e-156	433	100.00	358.1114	8BQF_A	8BQF_A
3	214	1	214	3.26e-156	432	100.00	358.0215	4X8M_A	4X8M_A
4	214	1	214	4.78e-156	432	100.00	357.6388	6S36_A	6S36_A
5	214	1	214	1.07e-155	431	99.53	356.8330	9R6U_A	9R6U_A
6	214	1	214	1.26e-155	431	99.53	356.6696	9R71_A	9R71_A

```
hits <- plot(blast)
```

* Possible cutoff values: 260 3
Yielding Nhits: 20 96

* Chosen cutoff value of: 260
Yielding Nhits: 20



```
hits$ pdb.id
```

```
[1] "1AKE_A" "8BQF_A" "4X8M_A" "6S36_A" "9R6U_A" "9R71_A" "8Q2B_A" "8RJ9_A"
[9] "6RZE_A" "4X8H_A" "3HPR_A" "1E4V_A" "5EJE_A" "1E4Y_A" "3X2S_A" "6HAP_A"
[17] "6HAM_A" "8PVW_A" "4K46_A" "4NP6_A"
```

```
# Download releated PDB files
#files <- get.pdb(hits$ pdb.id, path="pdbs", split=TRUE, gzip=TRUE)
```

```
# Align releated PDBs
#pdbs <- pdbaln(files, fit = TRUE, exefile="msa")
```

```
#pdbs
```

Quick interactive structure view

Principal Component Analysis

PCA of all this structural data (x, y and z atom coordinates):


```
# Perform PCA
#pc <- pca(pdbbs)
#plot(pc)
```

Interactive view of the PC1 captured structural differences:

```
# Create interactive view of the PC1 captured structural differences
#| eval: !expr knitr::is_html_output()
#view.pca(pc)
```